

Nataraja

End-to-End Architecture of the Cloud-Native LLM Orchestration Platform — how a request flows from edge to model to storage, and why every component is there.

9 request-path layers	3 cross-cutting planes	56 components explained
40+ tools & technologies	14 steps, edge → model	5 polyglot datastores


01 The request, end to end

Follow one job through the platform. A client submits it at the edge; the gateway accepts it instantly and returns an id; it is queued on the event backbone; a worker runs the agent loop against the model with resilience wrapped around every call; results stream back and persist across the polyglot stores. Each step below is one component in the path.

- 1 **Cloudflare CDN** — Serves cached content close to users and absorbs traffic before it hits our servers.
- 2 **AWS WAF** — Blocks malicious requests (SQL injection, bad bots) before they reach the app.
- 3 **Ingress controller** — Routes each incoming request to the correct internal service by host/path, terminating TLS.
- 4 **Next.js · React · TypeScript** — Where users submit jobs, watch progress live, and view queue / health / usage.
- 5 **FastAPI gateway** — Accepts a job and instantly returns 202 + a job id, so the user never waits.
- 6 **Keycloak · OAuth2/OIDC/JWT** — Authenticates every caller and issues tokens that say what they're allowed to do.
- 7 **Provider registry & selector** — Picks a healthy provider for each job (round-robin / weighted / least-loaded).
- 8 **Rate limiter** — Keeps our call rate within each provider's limits so we're never throttled or banned.
- 9 **Circuit breaker** — Stops hammering a provider that's already failing, then probes to see if it recovered.
- 10 **Kafka / Redpanda** — The queue that decouples the fast gateway from the slower workers; jobs wait here safely.
- 11 **Go workers** — Pulls jobs off Kafka and runs the agent to completion, many at once.
- 12 **Agent loop** — The core routine that lets the model break a goal into steps and solve it.
- 13 **Anthropic Claude API** — Does the actual reasoning and tool-use for every job (Messages API).
- 14 **PostgreSQL** — The source of truth for jobs, tenants, quotas and the billing ledger.

02 Architecture map

The full system as stacked layers — the request flows top to bottom; the three cross-cutting planes (platform, delivery, observability) support every layer above. Logos identify each product; symbols mark patterns.



LAYER 9 Storage plane · polyglot

right store for each data shape



PostgreSQL



MongoDB



Redis



pgvector / Qdrant



S3 / MinIO



ClickHouse

▽ SPANS EVERY LAYER ABOVE ▽

CROSS-CUTTING Platform & runtime

where everything runs



Kubernetes · EKS + GKE



Helm + Kustomize



Custom Operator + CRD



KEDA



HashiCorp Vault

CROSS-CUTTING Delivery & supply chain

how code reaches production



GitHub Actions + OIDC



Argo CD · GitOps



Argo Rollouts



Terraform / Crossplane



SBOM · cosign · SLSA

CROSS-CUTTING Observability & SRE

is it healthy? is it fast?



OpenTelemetry



Prometheus



Loki



Grafana



Pyroscope



SLOs + burn-rate alerts



Chaos Mesh



k6

03 Component reference

Every component in plain English — **what it is**, **its role** in the layer, **why we chose it**, and its trade-offs. Written so a non-technical reader can follow, and you can defend each choice.

Client & Edge



Cloudflare CDN

CLIENT & EDGE

What: A global content-delivery & edge network with 300+ locations worldwide.

Role: Serves cached content close to users and absorbs traffic before it hits our servers.

Why chosen: Massive free tier, built-in DDoS protection and WAF, and very low latency.

- + Global edge = fast for everyone
- + Built-in DDoS + bot protection
- + Cheap / generous free tier
- One more vendor to manage
- Cache invalidation needs care



AWS WAF

CLIENT & EDGE

What: A Web Application Firewall that inspects incoming HTTP requests.

Role: Blocks malicious requests (SQL injection, bad bots) before they reach the app.

Why chosen: Fully managed, ships with common attack rule-sets, integrates with our AWS load balancer.

- + Managed OWASP rule-sets
- + Blocks attacks at the edge
- Per-rule cost
- Needs tuning to avoid false blocks



Ingress controller

CLIENT & EDGE

What: A reverse proxy that is the single front door into the Kubernetes cluster.

Role: Routes each incoming request to the correct internal service by host/path, terminating TLS.

Why chosen: Nginx ingress is mature, battle-tested and flexible.

- + Mature & flexible routing
- + Central place for TLS & rules
- A choke point to scale & secure carefully



Route 53 · DNS

CLIENT & EDGE

What: Amazon's managed DNS service.

Role: Turns our domain name into the IP address of the load balancer, with health-based failover.

Why chosen: AWS-native, supports latency/geo routing and health checks for multi-region.

- + Health-checked failover
- + Latency & geo routing
- Ties DNS to AWS



cert-manager / ACM

CLIENT & EDGE

What: Automated TLS-certificate issuance and renewal.

Role: Gives every endpoint HTTPS and auto-renews certs so they never expire.

Why chosen: Automates free ACME/Let's Encrypt certs inside the cluster — no manual renewals.

- + Auto-renew = no expiry outages
- + Free certificates
- Depends on the cluster & DNS setup

Presentation



Next.js · React · TypeScript

PRESENTATION

What: The web dashboard, built with React + Next.js and typed with TypeScript.

Role: Where users submit jobs, watch progress live, and view queue / health / usage.

Why chosen: Server-side rendering, great developer experience, and types that catch bugs early.

- + Fast, SEO-friendly rendering
- + Types prevent whole classes of bugs
- + Huge ecosystem
- Build tooling is complex
- Frontend churn



SSE live stream

PRESENTATION

What: Server-Sent Events — a one-way live stream over plain HTTP.

Role: Pushes the agent's thinking to the browser token-by-token as it happens.

Why chosen: Simpler than WebSockets when you only need server→browser streaming.

- + Simple, HTTP-native, auto-reconnect
- One-way only
- Browser connection limits

API Gateway



FastAPI gateway

API GATEWAY

What: An async Python web framework serving as the platform's front door.

Role: Accepts a job and instantly returns 202 + a job id, so the user never waits.

Why chosen: Async-native (handles many requests at once), typed via Pydantic, auto-generates API docs.

- + Very fast & async
- + Typed request validation
- + Auto OpenAPI docs
- Requires async discipline
- Younger than Django/Flask



Idempotency + OCC

API GATEWAY

What: Idempotency keys plus optimistic concurrency control.

Role: Makes retries safe (no double-processing) and prevents two writers from clobbering data.

Why chosen: Clients and networks retry automatically — the system must never act twice on one request.

- + Safe retries
- + No lost updates
- Extra key storage & careful design



Keycloak · OAuth2/OIDC/JWT

API GATEWAY

What: An open-source identity provider speaking OAuth2 / OIDC / JWT.

Role: Authenticates every caller and issues tokens that say what they're allowed to do.

Why chosen: Full-featured, standards-based, self-hosted — no per-user SaaS fees.

- + Industry-standard auth
- + No per-user cost
- + Self-hosted control
- Heavier to operate yourself



Quotas + usage ledger

API GATEWAY

What: Per-tenant rate/usage limits plus an append-only usage record.

Role: Stops one customer starving the others and records usage for billing.

Why chosen: Multi-tenant fairness and accurate billing both need it.

- + Fair sharing
- + Billing-grade usage data
- Accounting adds complexity



Control plane



Provider registry & selector

CONTROL PLANE

What: A live registry of available LLM providers plus a routing strategy.

Role: Picks a healthy provider for each job (round-robin / weighted / least-loaded).

Why chosen: We need to route, weight and fail over between providers at runtime, not in code.

- + Flexible routing
- + Easy failover
- Must track provider health accurately



Experimentation / feature flags

CONTROL PLANE

What: An A/B-testing framework with feature flags.

Role: Rolls a new agent config or model out to a small % of traffic and measures results.

Why chosen: Lets us change the product with data and roll back instantly if a metric drops.

- + Safe, measurable rollouts
- + Instant rollback
- Flag sprawl
- Needs statistical rigor



Preview-env provisioner (Operator)

CONTROL PLANE

What: A custom Kubernetes Operator + CRD we write ourselves.

Role: Spins up and tears down isolated per-tenant environments on demand.

Why chosen: Automates what would otherwise be manual, error-prone ops work.

- + Self-service environments
- + Consistent & repeatable
- Real engineering effort to build

Resilience layer



Rate limiter

RESILIENCE LAYER

What: A token-bucket limiter backed by Redis + Lua.

Role: Keeps our call rate within each provider's limits so we're never throttled or banned.

Why chosen: Providers enforce hard rate limits; we must self-regulate across many workers.

- + Protects providers
- Distributed counting is tricky
- + Fair distributed limiting



Circuit breaker

RESILIENCE LAYER

What: A breaker with closed / open / half-open states.

Role: Stops hammering a provider that's already failing, then probes to see if it recovered.

Why chosen: Fail fast instead of piling requests onto a broken dependency.

- + Prevents cascading failures
- Thresholds need tuning
- + Auto-recovery probing



Retry + backoff + jitter

RESILIENCE LAYER

What: Smart retries that wait longer each time, with randomness (jitter).

Role: Recovers from brief, transient errors without all workers retrying in sync.

Why chosen: Networks blip; naive retries cause a stampede — jitter spreads them out.

- + Higher success rate
- Can amplify load if misconfigured
- + Avoids thundering herd



Fallback chain

RESILIENCE LAYER

What: An ordered list of backup providers.

Role: If the primary fails, quietly try the next one so the user still gets an answer.

Why chosen: Keeps the service useful during a partial outage.

- + Graceful degradation
- Can hide root causes — use carefully



Load shedding

RESILIENCE LAYER

What: Deliberately dropping excess requests under overload.

Role: Protects the system so it stays up for most users instead of collapsing for all.

Why chosen: Partial service beats total meltdown when demand exceeds capacity.

- + System stays alive under spikes
- Some requests are rejected

Event backbone



Kafka / Redpanda

EVENT BACKBONE

What: A distributed, durable event log (Redpanda is a drop-in Kafka).

Role: The queue that decouples the fast gateway from the slower workers; jobs wait here safely.

Why chosen: Durable, ordered, replayable and high-throughput — the backbone of event-driven systems.

- + Durable & replayable
- Operationally complex
- + Scales to huge throughput
- + Decouples producers/consumers



Schema Registry

EVENT BACKBONE

What: A registry of versioned event schemas.

Role: Lets producers and consumers evolve message formats without breaking each other.

Why chosen: In a big system, message shapes change — this makes that safe.

- + Safe schema evolution
- Extra infrastructure
- + Contract enforcement



Outbox pattern

EVENT BACKBONE

What: A transactional outbox table + relay.

Role: Publishes an event in the same DB transaction as the data write — no lost or duplicate events.

Why chosen: Guarantees the database and the event stream never disagree.

- + Strong consistency
- Extra table + relay process



Dead-letter queue

EVENT BACKBONE

What: A side queue for messages that repeatedly fail.

Role: Parks poison messages so one bad job doesn't block the whole stream.

Why chosen: Keeps the pipeline flowing while preserving failures for later inspection.

- + Isolates poison messages
- Needs monitoring & replay tooling



Debezium · CDC

EVENT BACKBONE

What: Change-Data-Capture that tails the database's change log.

Role: Streams every DB change into the analytics warehouse in real time.

Why chosen: Gets analytics data without risky dual-writes from the app.

- + Real-time, decoupled analytics
- Connectors add ops overhead



Worker fleet · Go



Go workers

WORKER FLEET · GO

What: The worker service that does the heavy lifting, written in Go.

Role: Pulls jobs off Kafka and runs the agent to completion, many at once.

Why chosen: Go's goroutines make massive concurrency cheap; binaries are tiny and fast.

- + Cheap concurrency
- A new language to learn
- + Fast & low memory
- + Tiny container images



Agent loop

WORKER FLEET · GO

What: The plan → act → observe reasoning cycle.

Role: The core routine that lets the model break a goal into steps and solve it.

Why chosen: This IS the product — an autonomous, tool-using agent.

- + Autonomous problem-solving
- Can loop or overspend if unbounded
- + Uses tools & data



Tool / function calling

WORKER FLEET · GO

What: Structured 'function calling' the model can invoke.

Role: Lets the agent take real actions (search, fetch, compute), not just talk.

Why chosen: Bridges the model's reasoning to the outside world.

- + Turns text into real actions
- Inputs must be validated for safety



Token budget

WORKER FLEET · GO

What: A cap on tokens (and thus cost/time) per job.

Role: Bounds how much a single job can spend before it must stop.

Why chosen: LLM calls cost real money — runaway agents must be contained.

- + Predictable cost
- May cut off long tasks
- + Prevents runaways



RAG retrieval

WORKER FLEET · GO

What: Retrieval-Augmented Generation over your own data.

Role: Finds relevant documents by meaning and feeds them to the model for grounded answers.

Why chosen: Cuts hallucinations and keeps answers current without retraining.

- + Grounded, up-to-date answers
- Only as good as the retrieval



gRPC + Protobuf

WORKER FLEET · GO

What: A fast, typed remote-procedure-call system.

Role: How internal services call each other efficiently with strict contracts.

Why chosen: Faster and stricter than REST for internal, high-volume service calls.

- + Fast binary protocol
- Not human-readable; extra tooling
- + Typed contracts
- + Streaming



Istio service mesh

WORKER FLEET · GO

What: A service mesh that sits beside every service.

Role: Adds mutual-TLS encryption and traffic control between services without app code.

Why chosen: Security and traffic-shaping handled by the platform, not each app.

- + mTLS everywhere
- Adds complexity & overhead
- + Canary & retries at mesh level
- + Free observability

LLM providers



Anthropic Claude API

LLM PROVIDERS

What: The large language model that powers the agents.

Role: Does the actual reasoning and tool-use for every job (Messages API).

Why chosen: Strong tool-use, long context and reliable reasoning — ideal for agents.

- + Excellent tool-use & reasoning
- + Long context window
- External cost & latency
- Provider rate limits

Storage plane · polyglot



PostgreSQL

STORAGE PLANE · POLYGLOT

What: A rock-solid relational (SQL) database.

Role: The source of truth for jobs, tenants, quotas and the billing ledger.

Why chosen: ACID guarantees, decades of maturity, plus JSON and extensions when needed.

- + Reliable & consistent
- + Extremely versatile
- Write-scaling needs planning



MongoDB

STORAGE PLANE · POLYGLOT

What: A document (NoSQL) database.

Role: Stores variable-shaped agent trajectories and payloads that don't fit tidy tables.

Why chosen: Flexible schema is perfect for deeply-nested, evolving agent data.

- + Flexible schema
- + Fast document reads
- Weaker multi-document transactions



Redis

STORAGE PLANE · POLYGLOT

What: An in-memory data store.

Role: Cache, hot state, and the counters behind rate-limiting and circuit-breaking.

Why chosen: Microsecond latency for the things that must be instant.

- + Blazing fast
- + Very versatile
- Memory-bound
- Durability trade-offs



pgvector / Qdrant

STORAGE PLANE · POLYGLOT

What: A vector database for embeddings.

Role: Stores the numeric 'meaning' of documents so RAG can find similar ones.

Why chosen: Enables semantic (meaning-based) search, not just keyword match.

- + Semantic similarity search
- Index tuning & memory cost



S3 / MinIO

STORAGE PLANE · POLYGLOT

What: Object storage (MinIO is S3-compatible, self-hosted).

Role: Holds large artifacts and long-term archives cheaply.

Why chosen: Effectively infinite, durable and inexpensive for blobs.

- + Cheap, durable, scalable
- Higher latency than a DB



ClickHouse

STORAGE PLANE · POLYGLOT

What: A columnar OLAP (analytics) database.

Role: The read-model / warehouse for fast analytics and dashboards.

Why chosen: Aggregates billions of rows in milliseconds — built for analytics.

- + Extremely fast analytics
- Not for transactional writes

Platform & runtime



Kubernetes · EKS + GKE

PLATFORM & RUNTIME

What: The container orchestrator that runs the whole fleet.

Role: Schedules, heals and scales every service across many machines.

Why chosen: The industry standard for running containers reliably at scale.

- + Self-healing & scaling
- + Portable across clouds
- Steep learning curve



Helm + Kustomize

PLATFORM & RUNTIME

What: Packaging and configuration tools for Kubernetes.

Role: Bundle each app's manifests and tailor them per environment.

Why chosen: Standard way to template and reuse deployment config.

- + Reusable, templated deploys
- Templating can get complex



Custom Operator + CRD

PLATFORM & RUNTIME

What: A controller we write that extends Kubernetes with our own resource type.

Role: Automates tenant creation and preview environments as native K8s objects.

Why chosen: Encodes our ops knowledge into the platform itself.

- + Automates complex ops
- Non-trivial to build correctly
- + Top-tier skill to show



KEDA

PLATFORM & RUNTIME

What: Event-driven autoscaler for Kubernetes.

Role: Adds or removes workers based on how many jobs are waiting in Kafka.

Why chosen: Scales on real backlog, not just CPU.

- + Scale on queue depth
- Another component to run
- + Scale to zero



HashiCorp Vault

PLATFORM & RUNTIME

What: A secrets manager and PKI.

Role: Issues short-lived database credentials and certificates on demand.

Why chosen: Dynamic, audited secrets beat long-lived passwords in config.

- + Dynamic, short-lived secrets
- Operationally involved
- + Full audit trail



Delivery & supply chain



GitHub Actions + OIDC

DELIVERY & SUPPLY CHAIN

What: The CI system that builds, tests and ships.

Role: Runs on every push and deploys to the cloud using short-lived OIDC tokens (no stored keys).

Why chosen: Keyless cloud auth removes the biggest secret-leak risk in CI.

- + Keyless, safer deploys
- YAML sprawl at scale
- + Tight GitHub integration



Argo CD - GitOps

DELIVERY & SUPPLY CHAIN

What: A GitOps continuous-delivery tool.

Role: Keeps the cluster exactly matching what's declared in Git, automatically.

Why chosen: Git becomes the single source of truth; drift self-corrects.

- + Declarative & auditable
- Everything must live in Git
- + Auto drift-correction



Argo Rollouts

DELIVERY & SUPPLY CHAIN

What: Progressive-delivery controller.

Role: Releases new versions gradually (canary / blue-green) and auto-rolls-back on bad metrics.

Why chosen: De-risks every deploy by testing on a slice of real traffic first.

- + Safe, gradual releases
- Adds release-config complexity
- + Metric-based auto-rollback



Terraform / Crossplane

DELIVERY & SUPPLY CHAIN

What: Infrastructure-as-Code tools.

Role: Define all cloud infrastructure (network, clusters, databases) as versioned code.

Why chosen: Reproducible, reviewable infrastructure instead of manual clicking.

- + Reproducible infra
- State management care needed
- + Peer-reviewable changes



SBOM - cosign - SLSA

DELIVERY & SUPPLY CHAIN

What: Software supply-chain security tooling.

Role: Lists everything in each image, signs it, and proves how it was built.

Why chosen: Lets us verify nothing was tampered with before it runs in production.

- + Tamper-evident builds
- Extra CI steps to maintain
- + Provenance & trust

Observability & SRE



OpenTelemetry

OBSERVABILITY & SRE

What: A vendor-neutral standard for traces, metrics and logs.

Role: Tags each request so we can follow it across every service (async, Kafka, gRPC).

Why chosen: One standard instrumentation instead of per-vendor agents.

- + Vendor-neutral
- + End-to-end request traces
- Instrumentation effort



Prometheus

OBSERVABILITY & SRE

What: A metrics database and alerting engine.

Role: Collects numbers (latency, error rate, queue depth) and fires alerts on rules.

Why chosen: The de-facto standard for cloud-native metrics.

- + Powerful queries
- + Strong alerting
- Long-term storage needs add-ons



Loki

OBSERVABILITY & SRE

What: A log aggregation system.

Role: Centralizes logs from every service so we can search them in one place.

Why chosen: Cheap, label-based logs that pair naturally with Grafana.

- + Cheap log storage
- + Grafana-native
- Not full-text like Elasticsearch



Grafana

OBSERVABILITY & SRE

What: The visualization layer.

Role: Dashboards that show the health of the entire system at a glance.

Why chosen: Single pane of glass over metrics, logs and traces.

- + One place for everything
- + Beautiful dashboards
- Dashboard maintenance



Pyroscope

OBSERVABILITY & SRE

What: Continuous profiling.

Role: Shows exactly which code lines burn CPU/memory in production.

Why chosen: Finds performance hot-spots that metrics alone can't.

- + Pinpoints hot code paths
- Small runtime overhead



SLOs + burn-rate alerts

OBSERVABILITY & SRE

What: Service Level Objectives with error budgets.

Role: Define 'good enough' reliability and alert only when we're burning the budget too fast.

Why chosen: Alerts on user-impact, not noise — the heart of SRE.

- + Meaningful, low-noise alerts
- + Aligns eng with users
- Requires honest target-setting



Chaos Mesh

OBSERVABILITY & SRE

What: A chaos-engineering tool.

Role: Deliberately injects failures (kill pods, add latency) to prove resilience.

Why chosen: You only know it's resilient if you've tested failure on purpose.

- + Proves resilience
- + Finds weak spots early
- Must be run carefully



k6

OBSERVABILITY & SRE

What: A load-testing tool.

Role: Simulates heavy traffic to find limits before real users do.

Why chosen: Scriptable, developer-friendly performance testing.

- + Realistic load tests
- + Scriptable in JS
- Test design takes effort